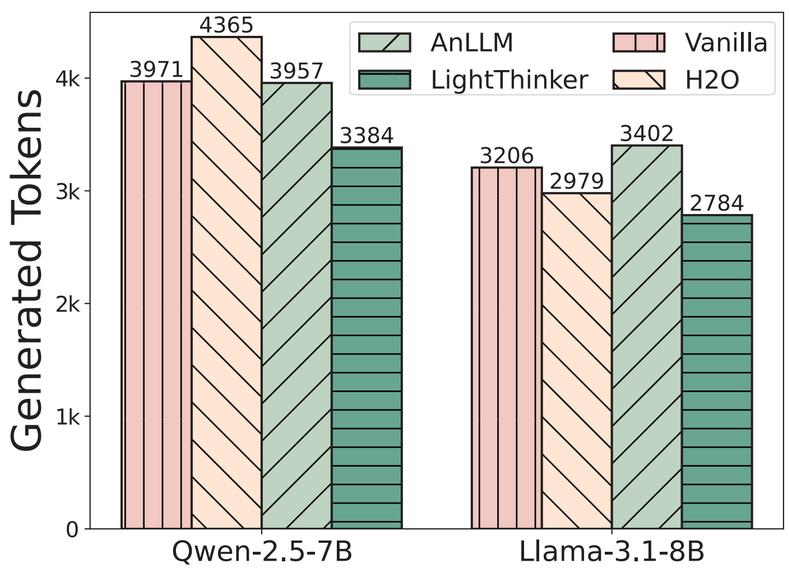

(a) Average number of generated tokens across the four datasets

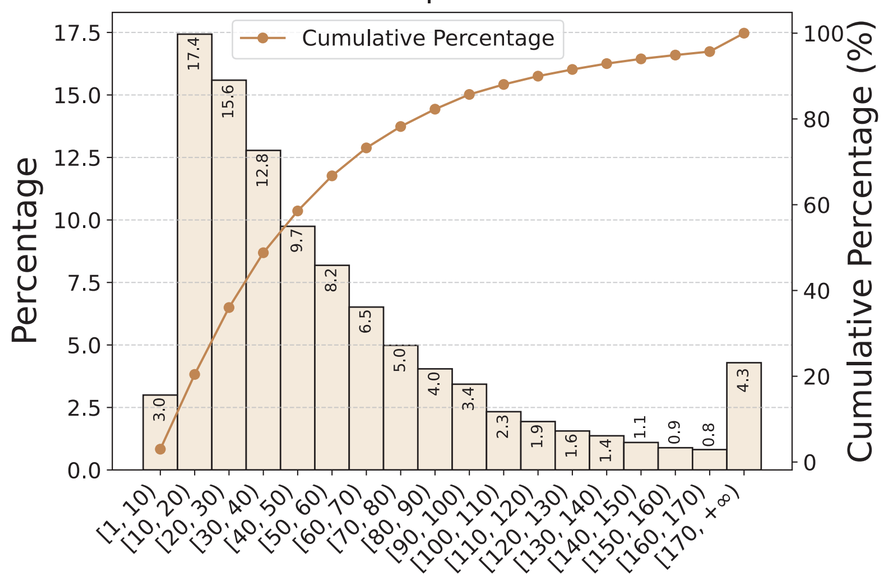

(b) Token compression frequency distribution for Qwen-2.5-7B on GPQA

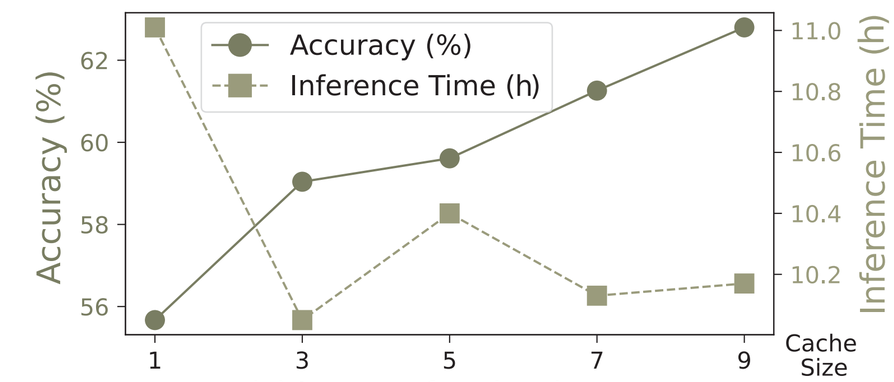

(e) Impact of cache size on accuracy and inference time

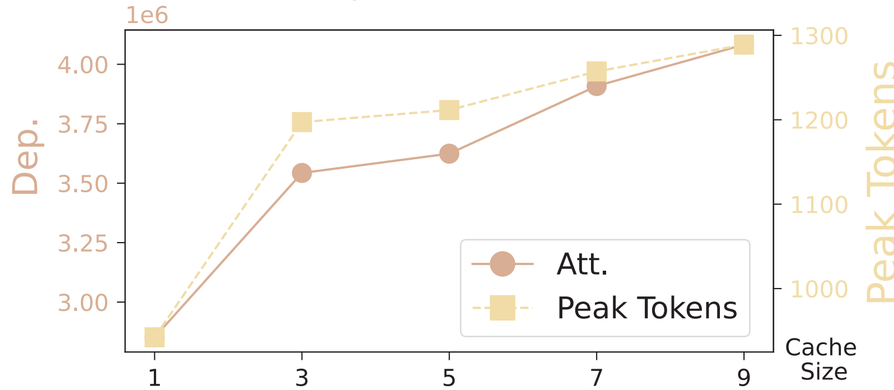

(f) Impact of cache size on Dep. and Peak tokens.

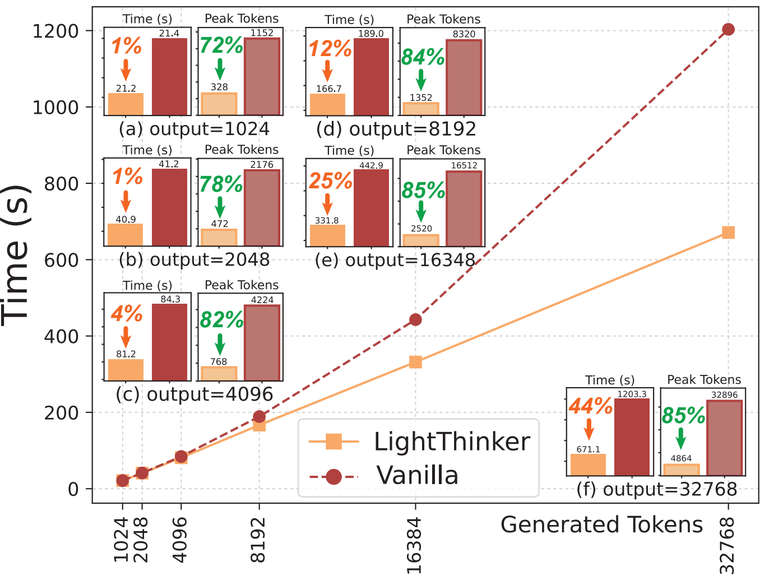

(c) Inference time vs. generated tokens Peak tokens vs. generated tokens

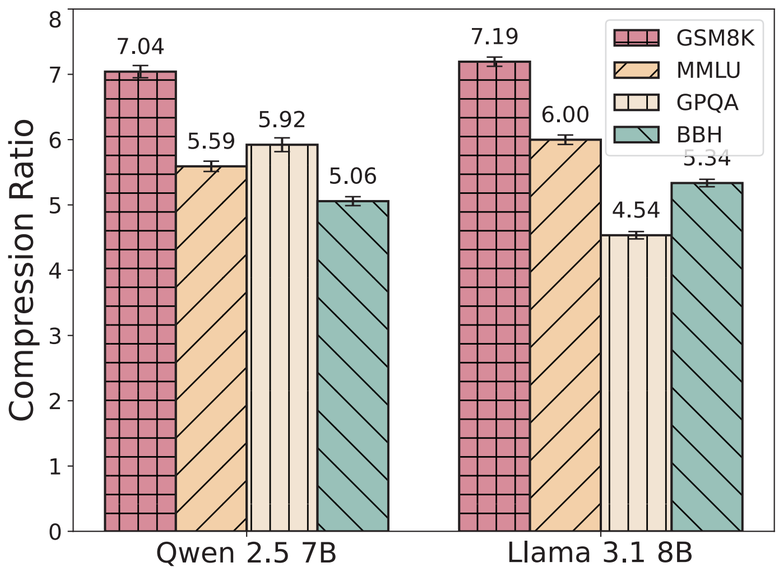

(d) Compression Ratios of LightThinker with 95% Confidence Intervals Across Multiple Datasets

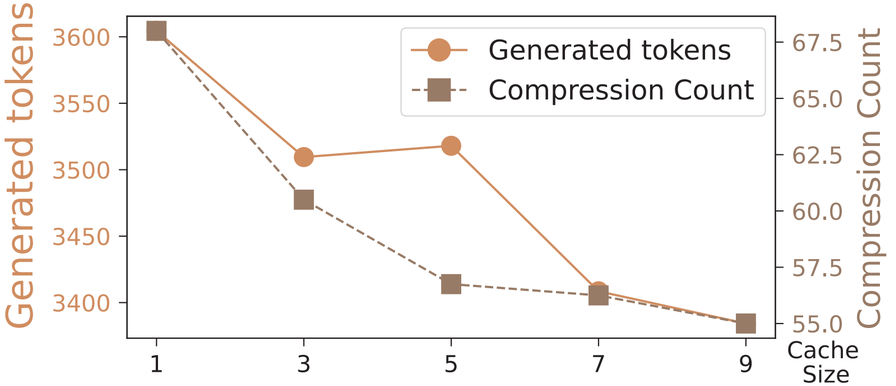

(g) Impact of cache size on generated token count and compression frequency